



From Prototype to Production

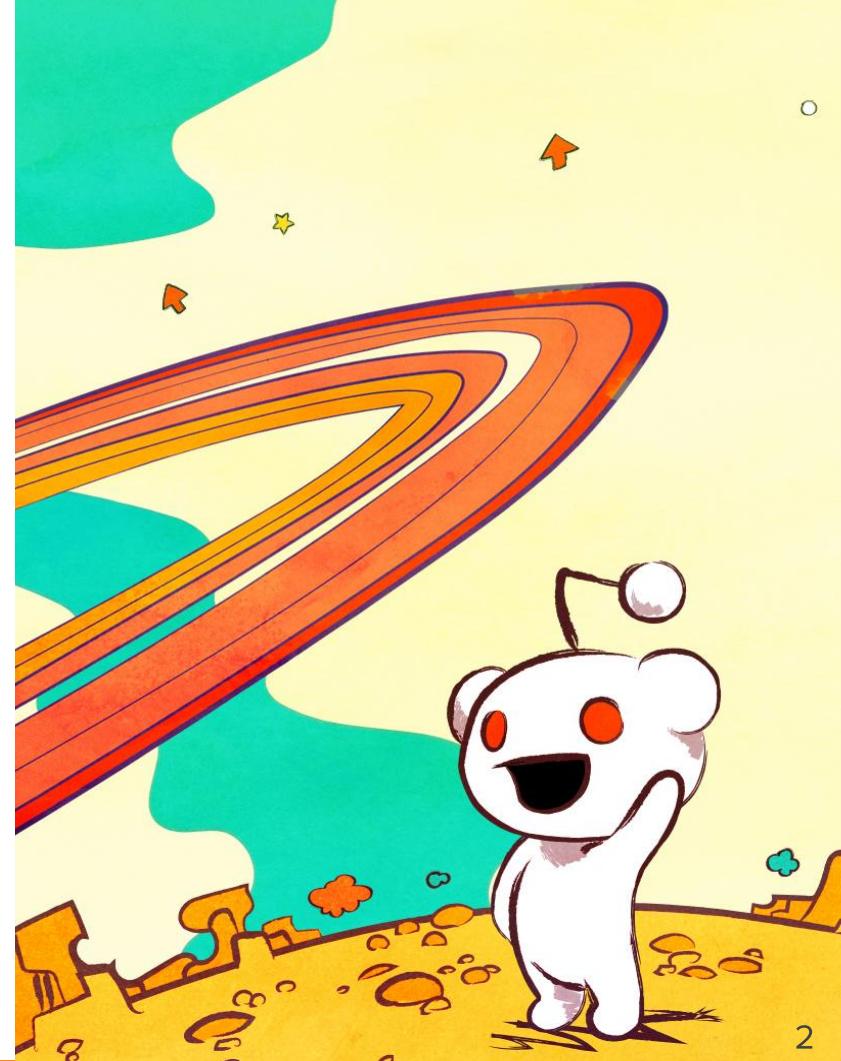
Lessons from Building and Scaling Reddit's Ad Serving Platform with Go

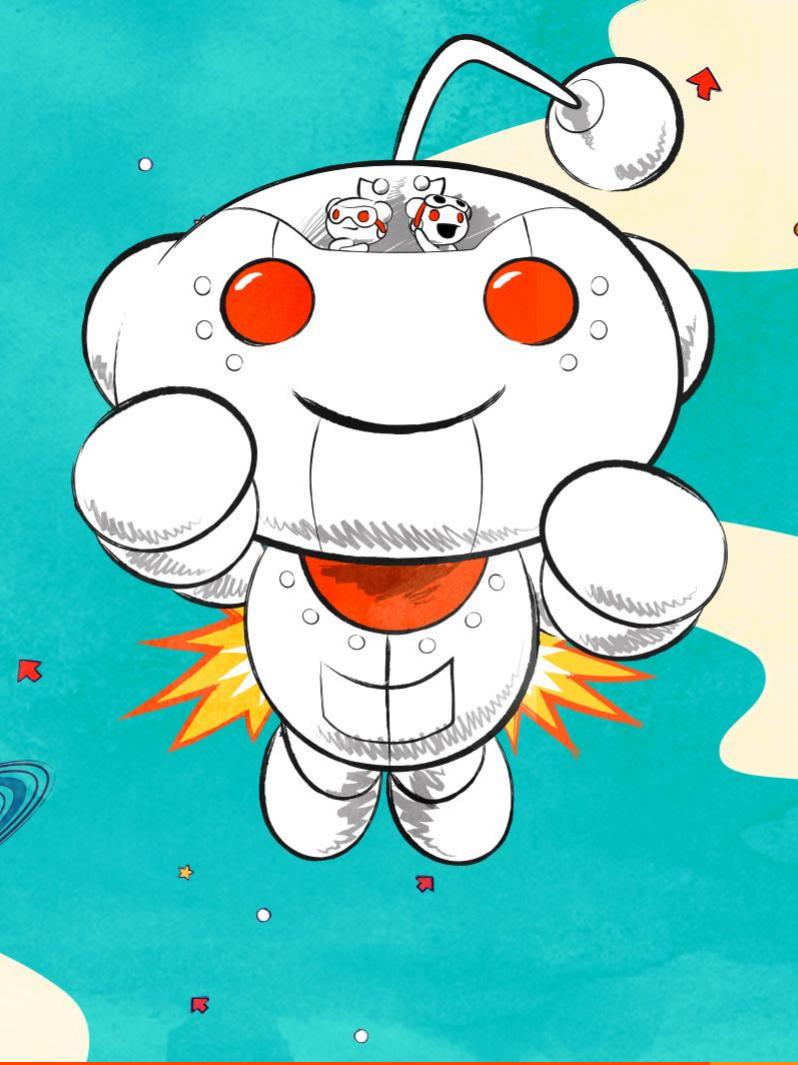
Deval Shah, Software Engineer, Reddit



Agenda

- Introduction to Reddit
- Ads Architecture Overview
- Our experience with Go
 - Lessons Learned
 - Challenges





What is Reddit?

Reddit is the **frontpage** of the internet

A social network where there are tens of thousands of **communities** around whatever **passions or interests** you might have

It's where people **converse** about the things that are most important to them



Reddit by the numbers

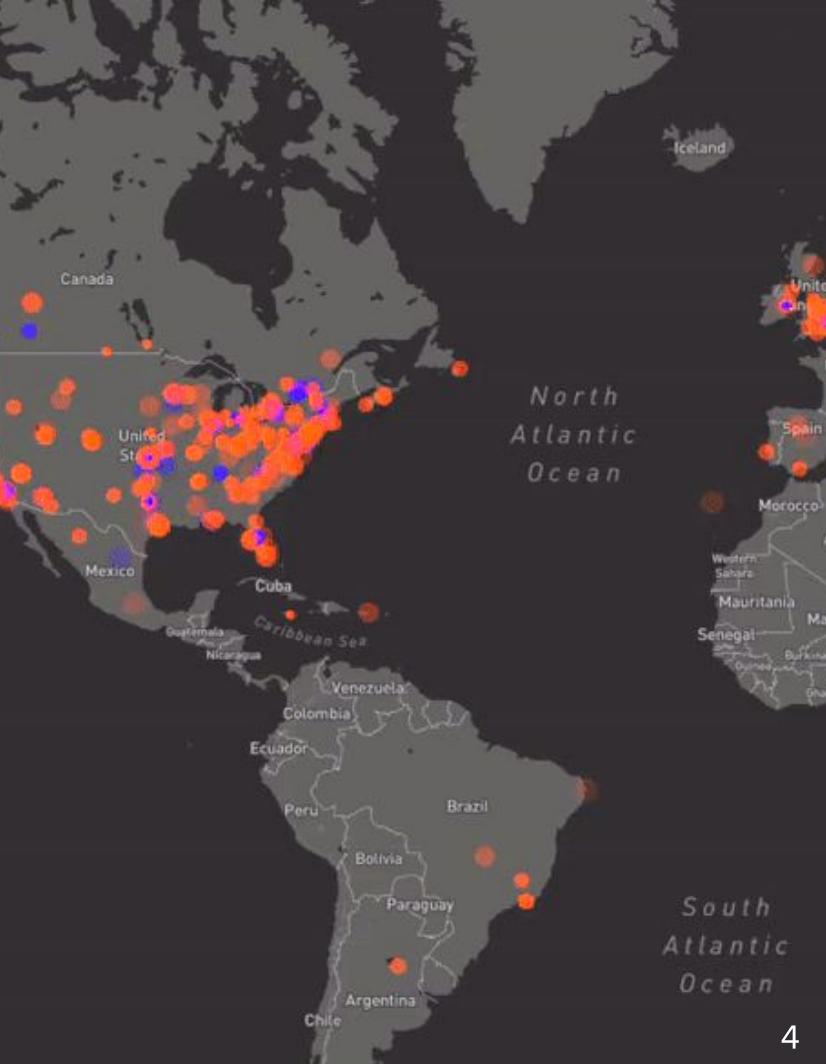
5th/18th Alexa Rank (US/World)

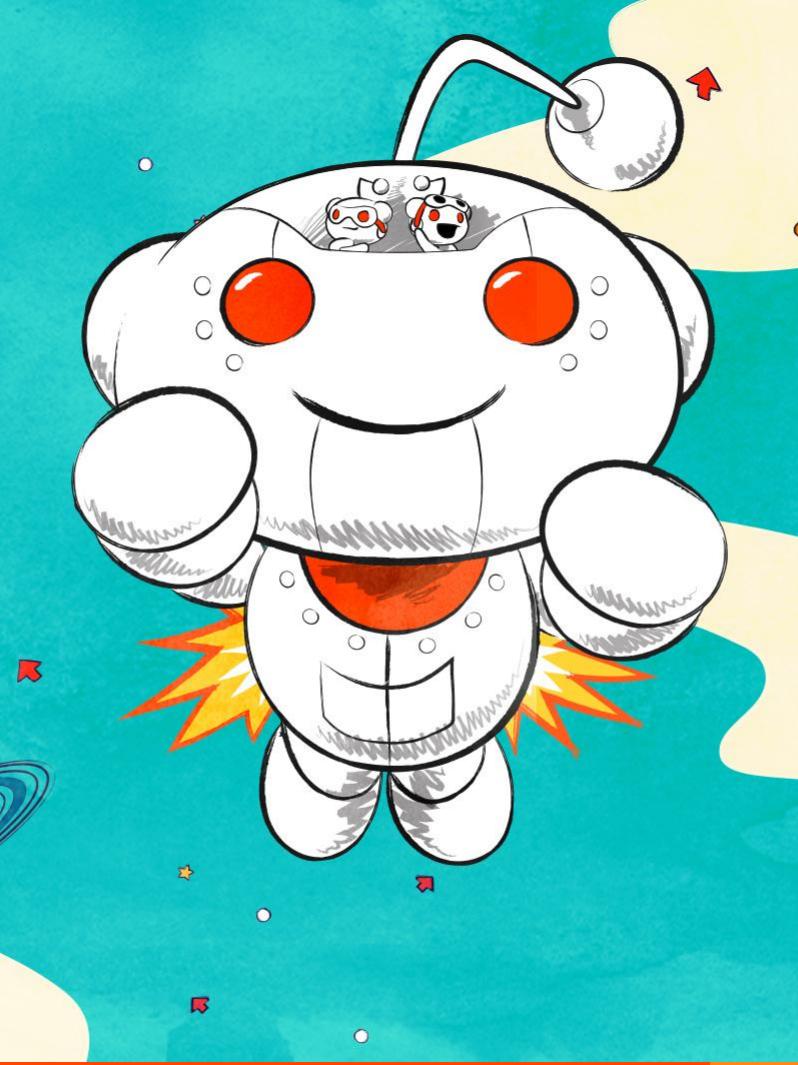
330M+ MAU

138K Active Communities

12M Posts per month

2B Votes per month





Ad Serving @ Reddit



Scale

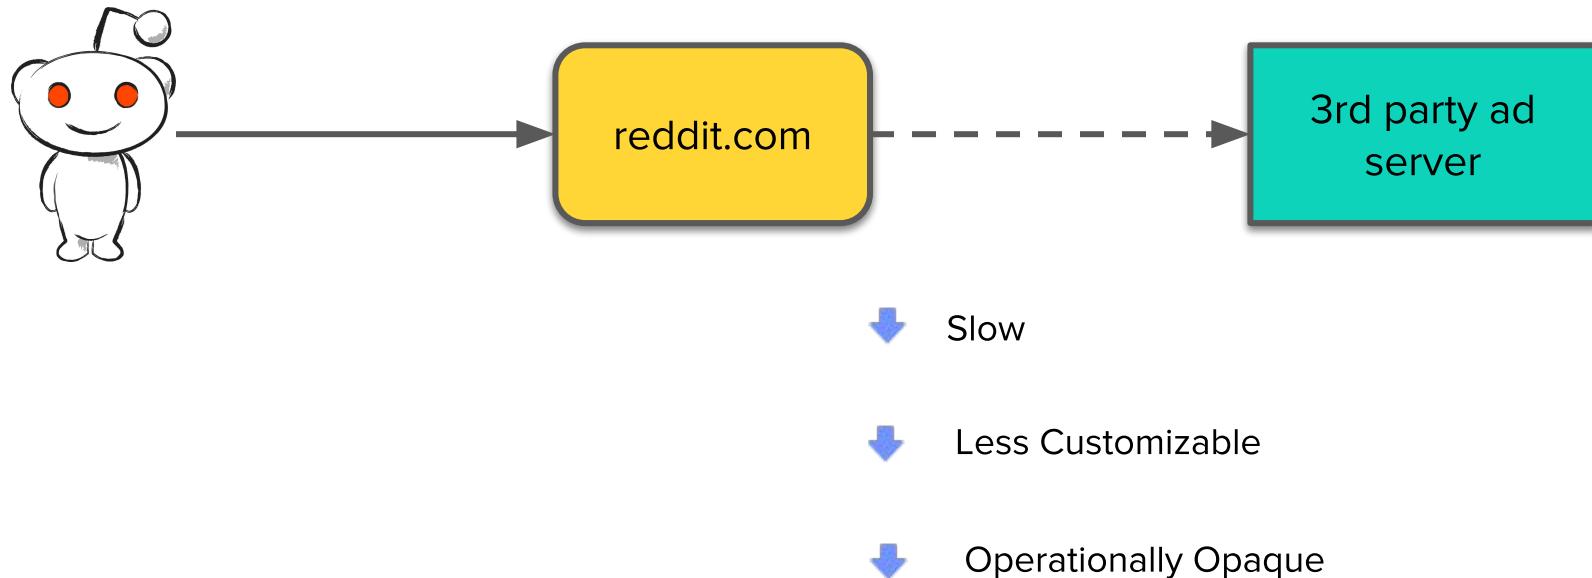
Speed

Auction

Pacing

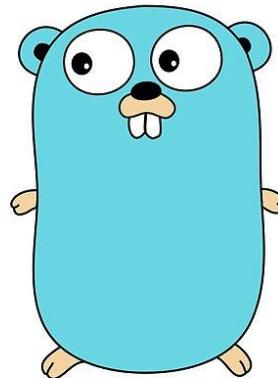


Ad Serving @ Reddit: Before

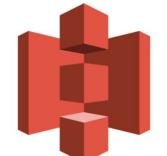




Ad Serving Infrastructure



Apache Thrift™



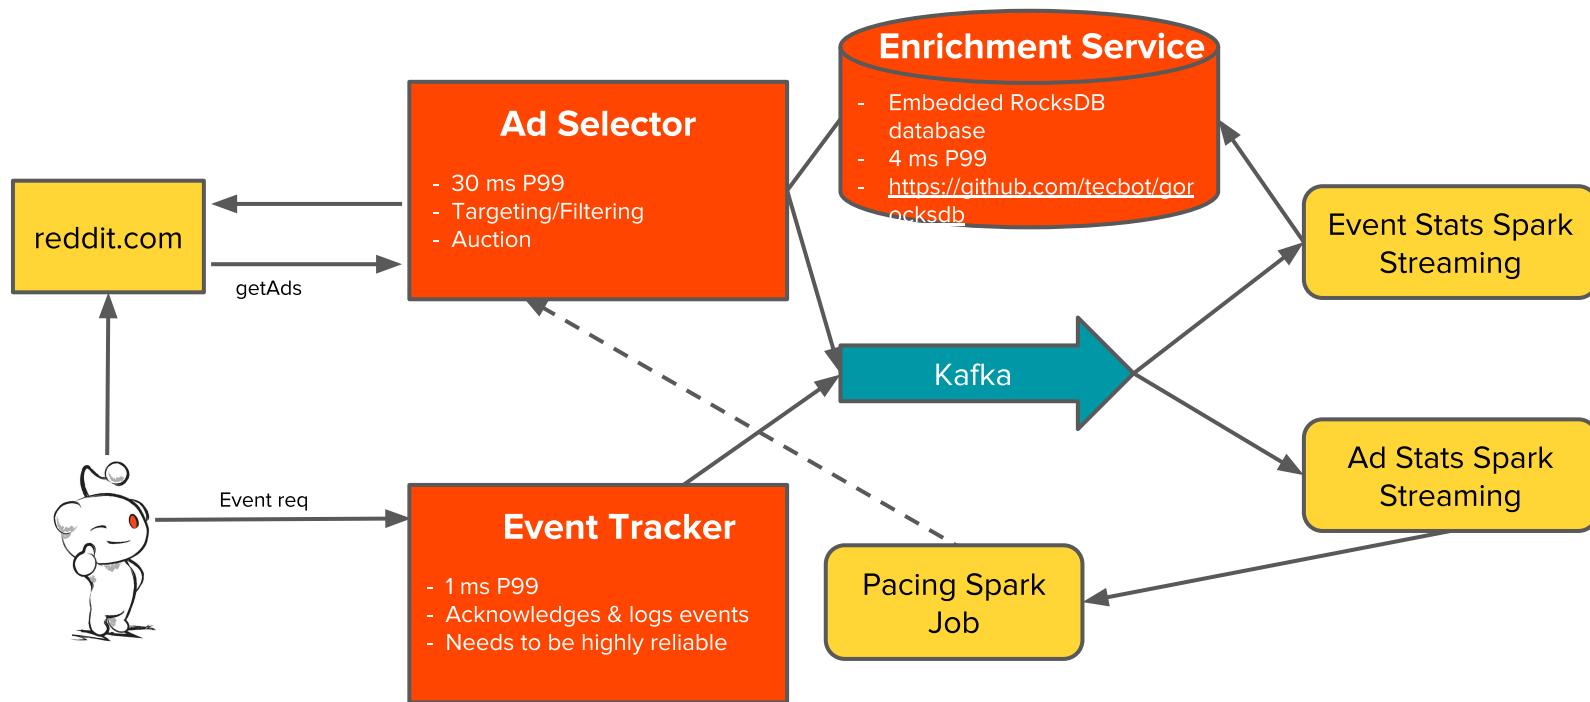
Amazon S3



Apache ZooKeeper™



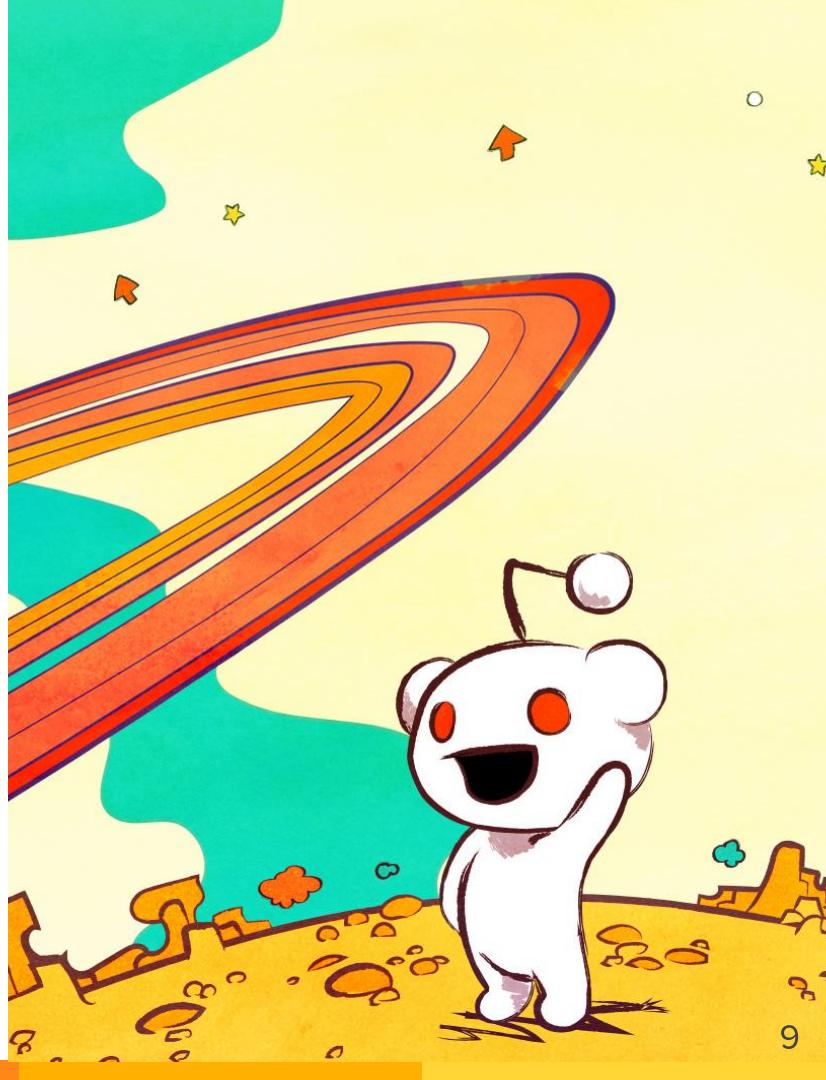
Ad Server Architecture





Other Go Tools & Services

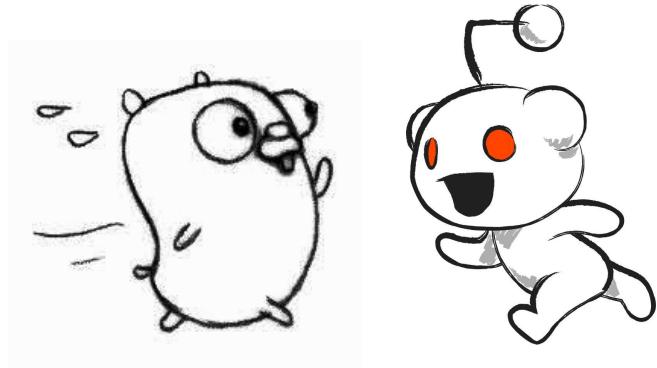
- Reporting Service
- Vault Administration Tool
- Ad Event Generation Service





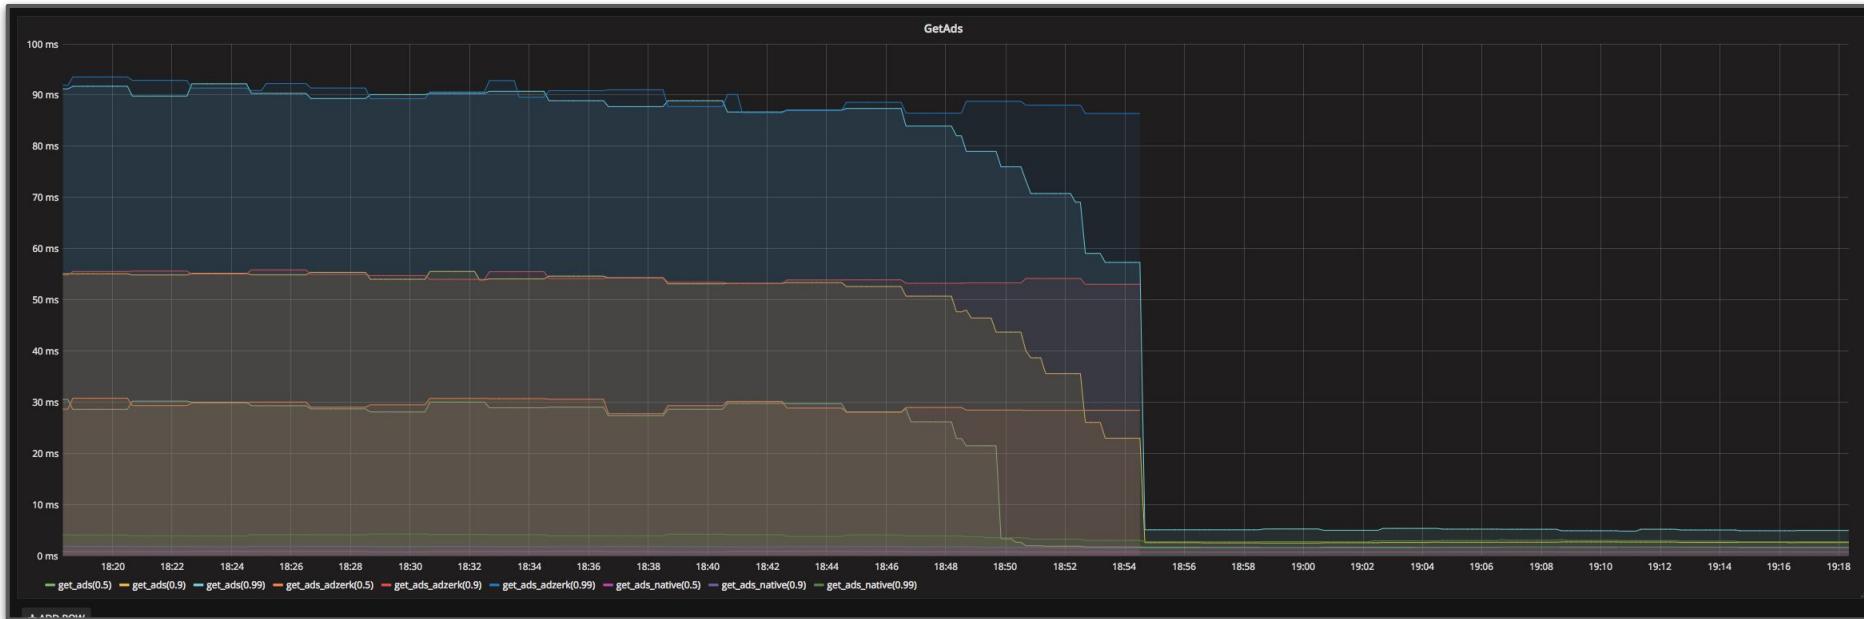
Our experience with Go

- Increased developer velocity
- Great performance out of the box
- Easy to focus on business logic
- Ad Serving Latency decreased drastically





Impact on ad serving latency



Lessons Learned





Problem 1

**How to build
production ready
microservices?**

- Initial prototype worked, but was not going to scale with developers
- Logging, metrics, etc. was all over the place
- Changing the transport layer was hard
- We needed repeatable patterns



Some Options

Go-Micro

“Pluggable RPC framework for distributed systems development”

Enforces gRPC

Great for beginners,
more opinionated than others

<https://github.com/micro/go-micro>

Gizmo

“A microservice toolkit that provides packages to put together server and pubsub daemons”

Supports gRPC, http

Good for pub-sub patterns

<https://github.com/NYTimes/gizmo>

Go-Kit

“A programming toolkit for building microservices in Go”

Supports gRPC, http, thrift

Flexible, not very opinionated

<https://gokit.io/>

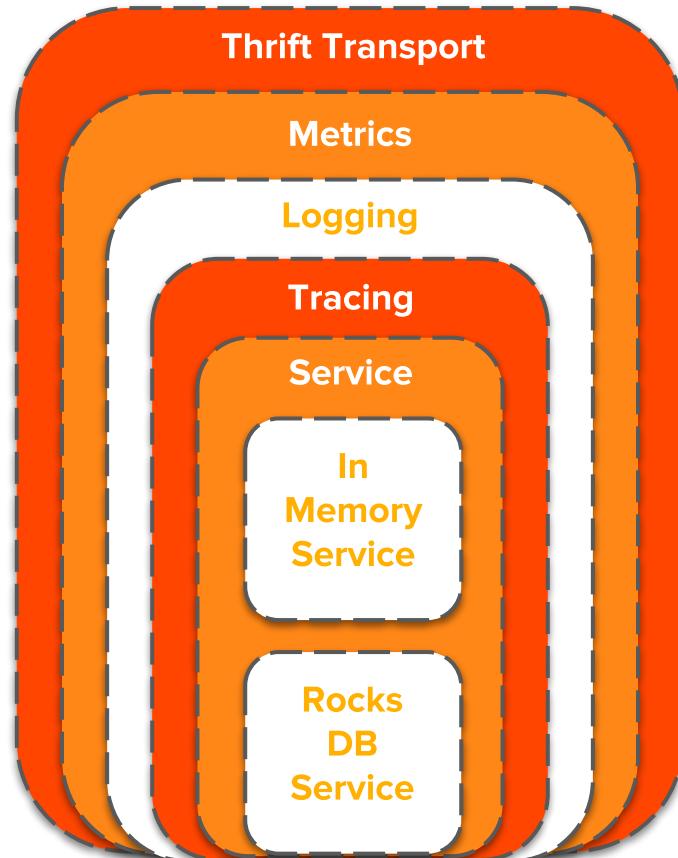


Why we picked Go-Kit

- Supports thrift
- Is flexible
- Has tools for logging, metrics, rate limiting, tracing, circuit breaking, etc.
- For more reasons, check out [Peter Bourgon's talk at GopherCon 2015](#)



Go-kit @ Reddit



1. Use a framework/toolkit



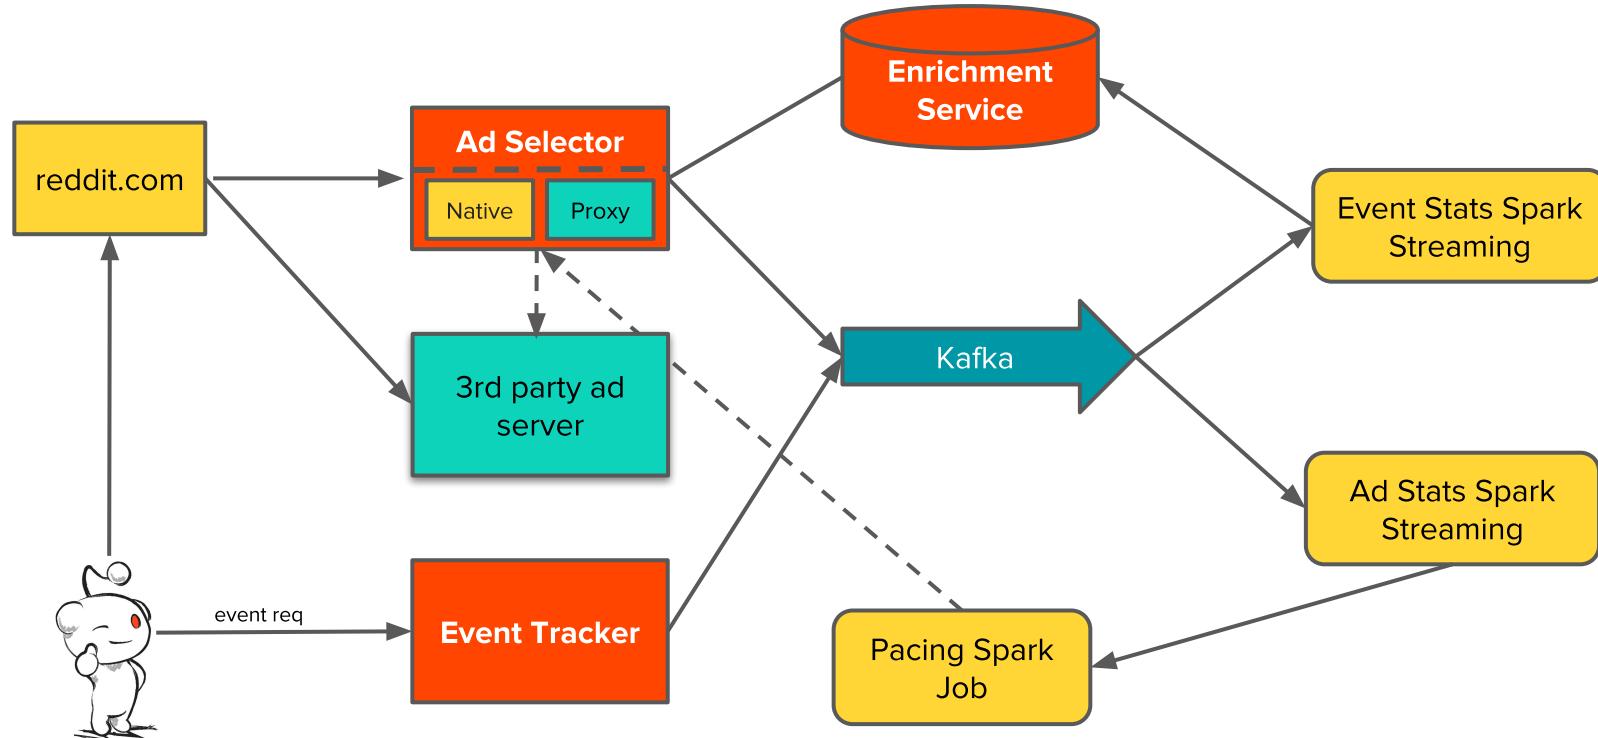
Problem 2

**How to roll out the
new system safely
& quickly?**

- Our goal was to roll out with minimal impact to
 - Reddit users
 - Advertisers
 - Other internal teams
- The third party ad server was a black box to us
 - Needed a way to iterate rapidly, learn and get better

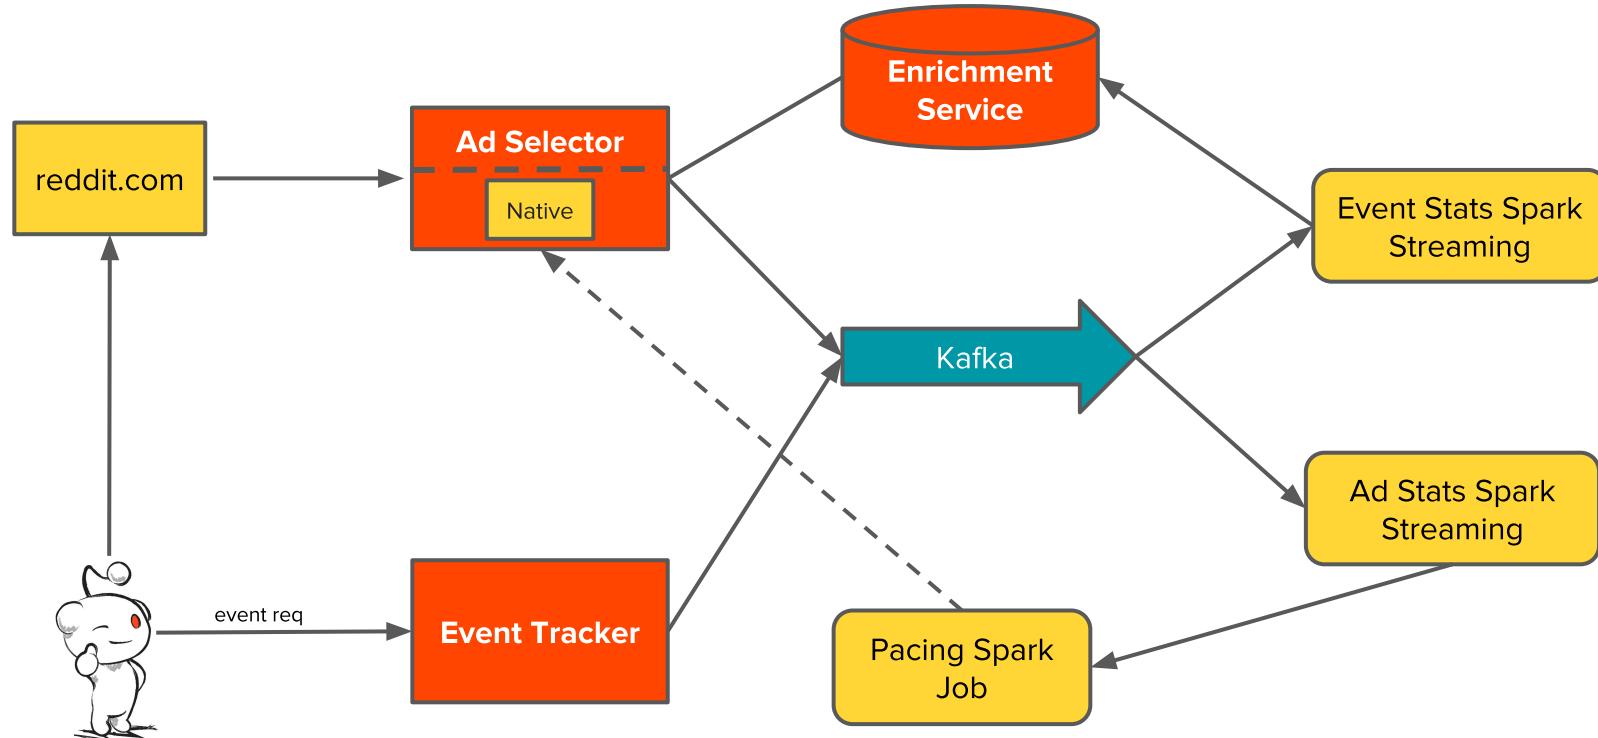


Changing Airplanes Mid-Flight (During Rollout)





Changing Airplanes Mid-Flight (After Rollout)





How did Go help with this?

- The Go Compiler is fast
- Cross platform compilation support
- Self-contained binary
- Strong concurrency primitives

2. Go makes rapid iteration easy & safe



Problem 3

How to debug latency issues?

- pprof is great if you know which service is having issues
- Distributed Tracing gives visibility across services



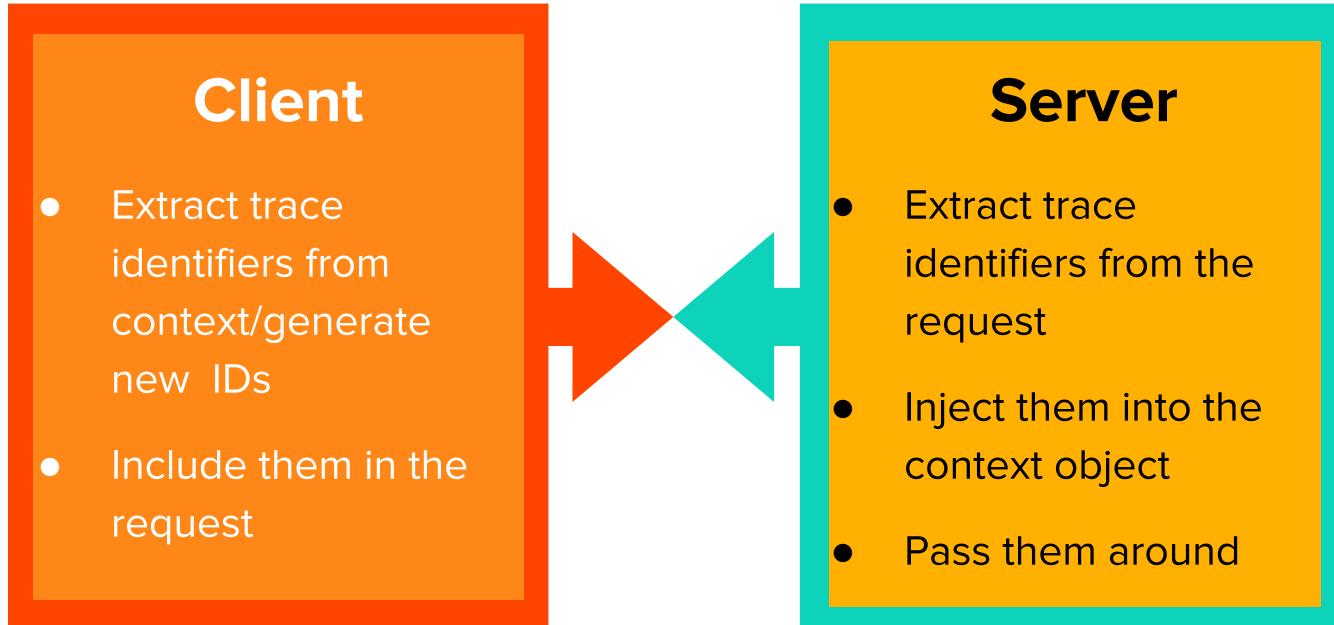
Why is Tracing useful?



- Identifies hotspots that cause high overall latency
- Helps find other errors/unexpected behavior



Tracing is usually easy



- See [go-kit's tracing](#) package for examples



Thrift Alternatives

Header

Context Object

FB Thrift

Apache Thrift





Solution: Add headers to Apache Thrift

- We added THeader to Apache Thrift
 - Context object is now supported
 - Headers can store trace identifiers
- <https://github.com/devalshah88/thrift>



Client Wrapper

```
// Call injects trace identifiers from ctx into the thrift headers
func (p *TracingThriftClient) Call(ctx context.Context, method string, args, result thrift.TStruct) error {

    headerProto, ok := p.oprot.(*thrift.THeaderProtocol)
    if ok {
        // Get trace info from ctx, start new span if one doesnt exist
        var clientSpan opentracing.Span
        if parentSpan := opentracing.SpanFromContext(ctx); parentSpan != nil {
            clientSpan = p.tracer.StartSpan(
                method,
                opentracing.ChildOf(parentSpan.Context()),
            )
        } else {
            clientSpan = p.tracer.StartSpan(method)
        }
        defer clientSpan.Finish()
        ext.SpanKindRPCClient.Set(clientSpan)

        // Inject trace identifiers into the header protocol
        if err := p.tracer.Inject(clientSpan.Context(),
            opentracing.TextMap,
            THeaderReaderWriter{headerProto}); err != nil {
            p.logger.Log("err", err)
        }
    }

    return p.Client.Call(ctx, method, args, result)
}
```



Server Wrapper

```
// Extract headers into context object
headerProto, ok := iprot.( *thrift.THeaderProtocol )
if ok {
    var span opentracing.Span
    wireContext, err := w.tracer.Extract(opentracing.TextMap, THeaderReaderWriter{headerProto})
    if err != nil && err != opentracing.ErrSpanContextNotFound {
        w.logger.Log("err", err)
    } else {
        span = w.tracer.StartSpan(w.name, ext.RPCSvServerOption(wireContext))
    }
    ctx = opentracing.ContextWithSpan(ctx, span)
}
```

<https://github.com/devalshah88/thrift-tracing>

Success!

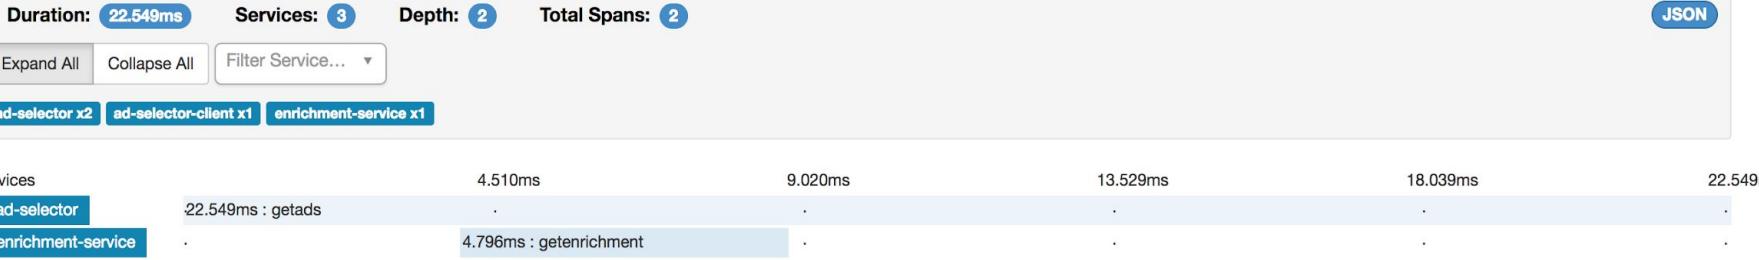
Zipkin Investigate system behavior

Find a trace

View Saved Trace

Dependencies

Go to trace



3. Distributed tracing with Thrift & Go is hard



Problem 4

How to handle slowness/timeouts?

- Don't make users wait for too long
- Don't waste resources doing unnecessary work



Use context to enforce timeouts within a service

```
d := time.Now().Add(20 * time.Millisecond)
ctx, cancel := context.WithDeadline(context.Background(), d)

defer cancel()

ch := make(chan EnrichmentData, 1)

go getEnrichmentData(ctx, ch)

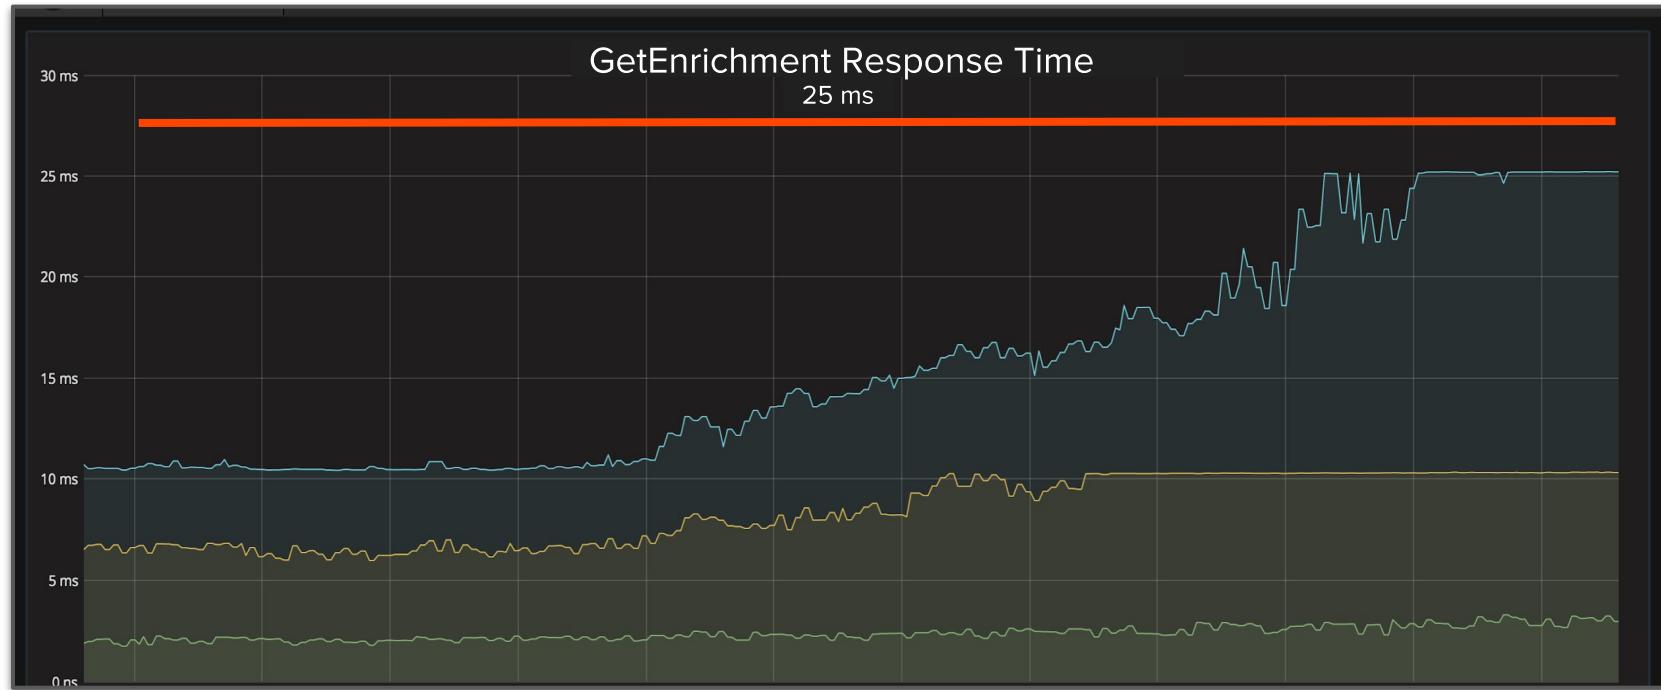
select {

case data := <-ch:
    return &data, nil

case <-ctx.Done():
    return nil, ctx.Err()
}
```

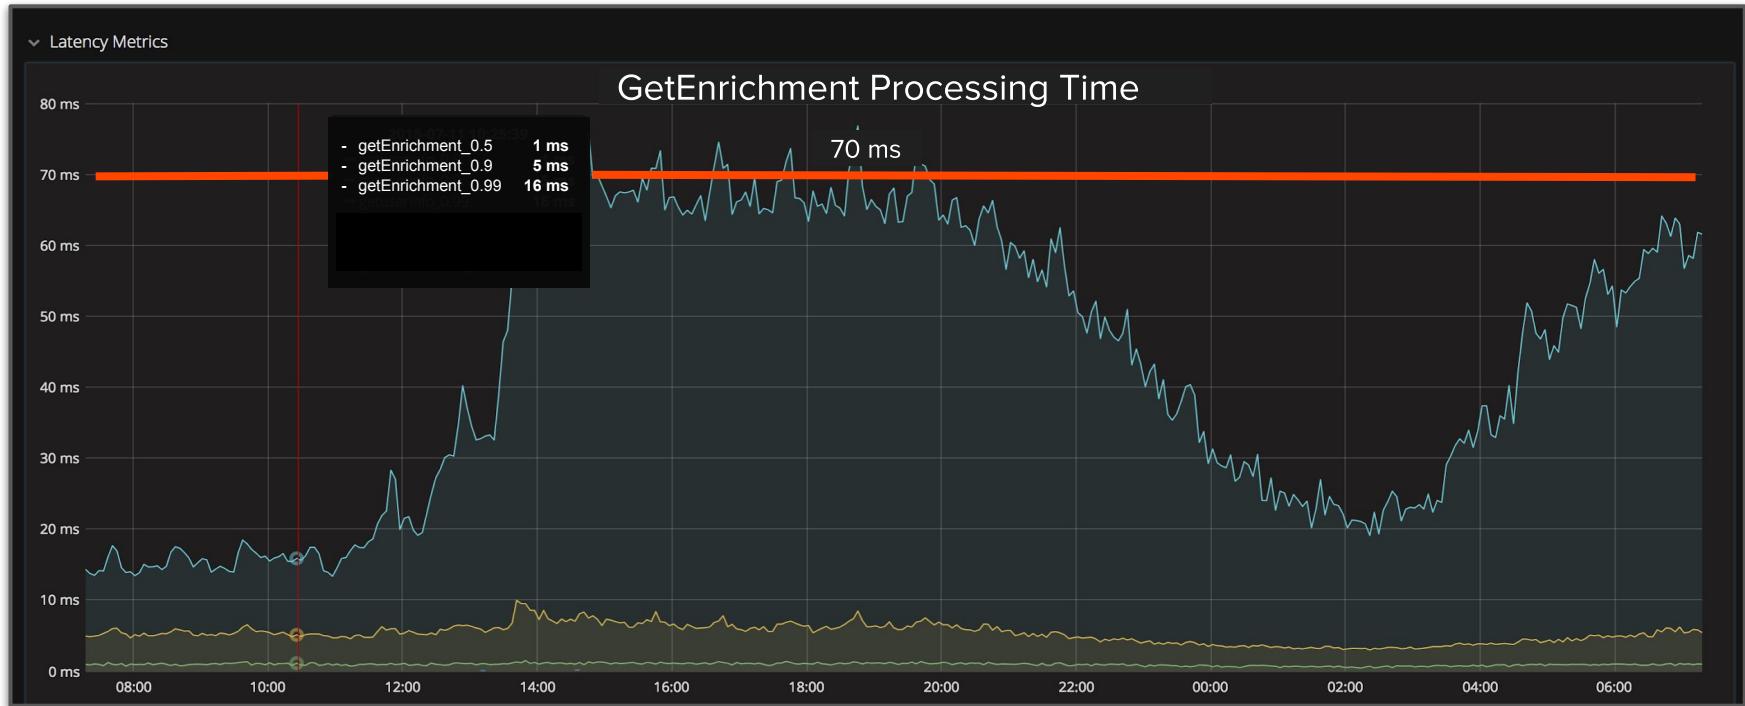


This is good





But not enough





Propagating deadlines with HTTP

```
uri := "https://example.com"
req, err := http.NewRequest("GET", uri, nil)
if err != nil {
    log.Fatalf("NewRequest() failed with '%s'\n", err)
}

ctx, cancel := context.WithTimeout(ctx, time.Millisecond*100)
defer cancel()

req = req.WithContext(ctx)

resp, err := http.DefaultClient.Do(req)

if err != nil {
    log.Fatalf("failed with:\n'%s'\n", err)
}
defer resp.Body.Close()
body, err := ioutil.ReadAll(resp.Body)
return body
```



Thrift makes this hard

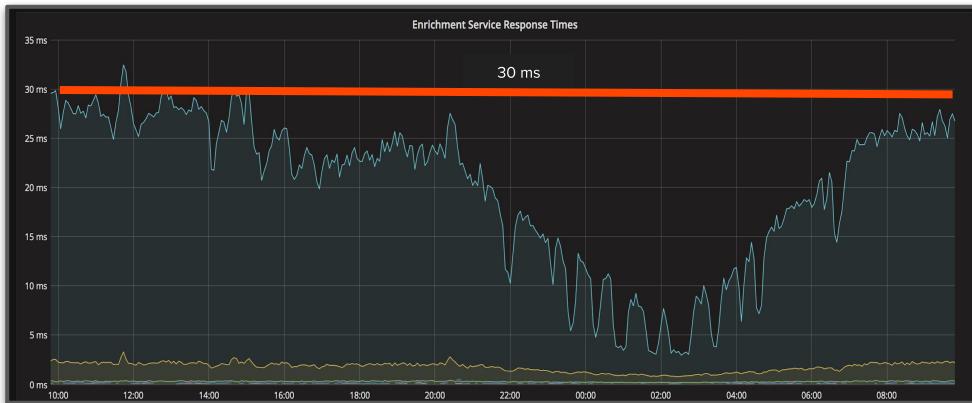
```
func (p *TSimpleServer2) AcceptLoop() error {
    for {
        client, err := p.serverTransport.Accept()
        if err != nil {
            select {
            case <-p.quit:
                return nil
            default:
            }
            return err
        }
        if client != nil {
            p.Add(1)
            go func() {
                if err := p.processRequests(client); err != nil {
                    log.Println("error processing request:", err)
                }
            }()
        }
    }
}
```

- Context object not used for cancellations/timeouts
- Goroutine leak



But it can be done

- Add deadline to the request payload
 - Client: Include deadline in request
 - Server: Inject deadline into context and use it
- Pass deadline as a thrift header
 - Similar to how we pass trace identifiers



4. Use Deadlines within & across services



Problem 5

How to ensure new features don't degrade performance?

- Rapid iteration & complex business logic could lead to performance issues
- Needed processes/tools to ensure we could move fast without violating our latency SLA



Load Testing using Bender

```
func main() {
    flag.Parse()

    intervals := bender.ExponentialIntervalGenerator(float64(*qps))
    exec := bthrift.NewThriftRequestExec(
        thrift.NewTHeaderTransportFactory(thrift.NewTBufferedTransportFactory(8192)),
        GetAdsExecutor, 10*time.Second, *addr)
    recorder := make(chan interface{}, 10000)
    requests := GetAdsRequests(*count)
    bender.LoadTestThroughput(intervals, requests, exec, recorder)

    l := log.New(os.Stdout, "", log.LstdFlags)
    h := hist.NewHistogram(60000, int(time.Microsecond))
    bender.Record(recorder, bender.NewLoggingRecorder(l), bender.NewHistogramRecorder(h))
    fmt.Println(h)
}
```



Load Testing

```
Percentiles:  
  Min:      536  
  Median:   3707  
  90th:     5796  
  99th:     6202  
  99.9th:   10356  
  99.99th:  30503  
  Max:     60000  
  
Stats:  
  Average: 3873.150700  
  Total requests: 10000  
  Elapsed Time (sec): 19.8343  
  Average QPS: 504.18  
  Errors: 0  
  Percent errors: 0.00
```

<https://github.com/pinterest/bender>



Benchmarking

```
func BenchmarkFetchAds(b *testing.B) {
    // Setup
    selector := Selector{&testDomain, &testSecret, &blankCampaignID}
    request := makeRequest(())
    selector.UpdateAdAvailability(testCtx, makeAdEngagements(1000, 8))

    // Run Benchmark
    for n := 0; n < b.N; n++ {
        selector.FetchNativeAds(context.Background(), request)
    }
}
```



Benchmarking

```
Running tool: /usr/local/go/bin/go test -benchmem -run=^$ ad_selector/native -bench ^BenchmarkFetchAds$  
  
goos: darwin  
goarch: amd64  
pkg: ad_selector/native  
BenchmarkFetchAds-8          10000    112662 ns/op      58762 B/op      605 allocs/op  
PASS  
ok    ad_selector/native 1.179s  
Success: Benchmarks passed.
```

5. Load testing & Benchmarking is easy. Do it!



Let's Recap

1. Use a framework/toolkit
2. Go makes rapid iteration safe & easy
3. Distributed tracing with thrift & Go is hard
4. Use Deadlines within & across services
5. Load testing & Benchmarking is easy. Do it!



Conclusion

- Go has been instrumental in helping reddit build and scale the new ad serving platform
- We shared 5 key takeaways from our experience
- Try to use at least one of them in your next Go project!



Thanks!
Deval Shah
twitter.com/devalshah

